

Data Structures

Andrew Rosen

Contents

I	Preliminaries	9
1	Introduction	11
1.1	What is a Data Structures Course	11
1.2	Why This Book?	11
1.2.1	Where Does This Book Fit Into a Computer Science Curriculum	11
1.2.2	What Are My Base Assumptions about the Reader? . . .	11
1.3	To The Instructor	12
1.3.1	How to Use	12
1.4	To The Student	12
1.4.1	How to use	12
2	Functions and How They Work	13
2.1	The Runtime Stack	13
2.2	Argument vs Parameter	13
2.2.1	Does anyone actually care?	13
2.3	Passing Arguments	13
2.3.1	How it Works in Java	13
2.3.2	How it works in Python	13
3	The Array	15
3.1	Why Arrays	15
3.2	Java and Arrays	15
3.3	Python and Arrays	16
3.3.1	Cool Ways to Build an Array in Python	16
3.4	How an Array Works	16
3.4.1	Operations	16
3.4.2	Array Internals and the Memory Formula	16
3.5	Common Array Algorithms	16
3.5.1	Finding Values in an Array	16
3.5.2	Limitations	16
4	Analyzing Algorithms	17
4.0.1	Cost	17
4.1	Big O Notation	17
4.1.1	Space Complexity	18
4.2	Examples with Arrays	18
4.2.1	Selection Sort	18

4.2.2	Bubble Sort	18
4.2.3	Insertion Sort	18
4.2.4	Other Sorting Algorithms	18
4.3	The Formal Mathematics of Big O Notation	18
4.4	Other Notations	18
4.5	When To Ignore Costs	18
II	Lists	19
5	Array Lists	21
5.1	What is a List?	21
5.1.1	A note on terminology	22
5.1.2	Lists in Java	22
5.2	Generics	22
5.2.1	What are they?	22
5.2.2	But Why?	22
5.3	List Operations	22
5.3.1	Size	22
5.3.2	Add	23
5.3.3	Remove	23
5.3.4	Get	23
5.3.5	Set	23
5.4	ArrayLists	23
5.5	Example Algorithms	24
5.6	Building an ArrayList	24
5.6.1	Caveats	24
5.6.2	Instance Variables	25
5.6.3	Constructor	25
5.6.4	Size	27
5.6.5	Add	27
5.6.6	toString and <code>_str_</code>	28
5.6.7	Get and Set	28
5.6.8	Remove	30
5.7	Analysis	30
5.8	More Restrictive or Permissive Generics	30
5.9	A Few More Useful Methods	30
5.9.1	Constructors	30
5.9.2	Trim	30
5.9.3	Adding Multiple Items in One Invocation	30
6	Linked Lists	33
6.1	Connecting Nodes into a list.	34
6.2	Building a Singly LinkedList	34
6.2.1	The Node	34
6.2.2	Instance Variables and Constructor	35
6.2.3	Adding	35
6.3	Get and Set	39
6.3.1	Get	39
6.3.2	Set	39

<i>CONTENTS</i>	5
-----------------	---

6.4 Remove	40
6.5 Analysis	40
6.5.1 Some Algorithms Play Better	40
6.6 Potential Project/Practice/Labs	40
6.7 Source Code	40

7 Stacks 43

7.1 Stack Operations	43
7.2 Building a Stack	43
7.3 Built-in Stacks	43
7.4 Solving Problems with A Stack	43
7.5 Mazes - Stacks and Backtracking	43
7.6 Discrete Finite Automata	43

8 Queues 45

8.1 Linked Based Implementation	45
8.2 Array Based Implementation	45

III Recursion 47

9 Recursion 49

9.1 Introduction	49
9.2 Recursive Mathematics	49
9.2.1 Fibonacci	49
9.3 Redoing Things With Recursion	49
9.3.1 Printing Recursively	49
9.3.2 Recursive Linear Search	49
9.4 Binary Search	49
9.4.1 Runtime Analysis	49
9.5 Recursive Backtracking	50
9.5.1 Mazes Again	50
9.5.2 The Eight Queens Puzzle	50
9.5.3 Additional Problems left to the Reader	50
9.6 Recursive Combinations	50
9.7 Recursion and Puzzles	50
9.8 Recursion and Art	50
9.9 Recursion and Nature	50

10 Trees 51

10.1 The Parts of a Tree	51
10.1.1 Where the Recursion comes in	52
10.2 Binary Search Trees	52
10.3 Building a Binary Search Tree	52
10.3.1 The Code Outline	52
10.3.2 Add	53
10.3.3 Contains	53
10.3.4 Delete	53

11 Heaps	55
11.1 Priority Queues	55
11.2 Removing From other locations	55
12 Sorting	57
12.1 Quadratic-Time Algorithms	57
12.1.1 Bubble Sort	57
12.1.2 Selection Sort	57
12.1.3 Insertion Sort	57
12.2 Log-Linear Sorting Algorithms	57
12.2.1 Tree Sort	57
12.2.2 Heap Sort	57
12.2.3 Heapify	58
12.2.4 Quick Sort	58
12.2.5 Merge Sort	58
12.3 Unique Sorting Algorithms	58
12.3.1 Shell Sort	58
12.3.2 Radix Sort	58
12.4 State of the Art Sorting Algorithms	58
12.4.1 Tim Sort	58
12.4.2 Quick Sort	58
12.5 But What if We Add More Computers: Parallelization and Dis-	
tributed Algorithms	58
12.6 Further Reading	58
12.6.1 Pedagogical Sorting Algorithms	58
IV Hashing	59
13 Sets	61
13.1 Operations	61
13.1.1 Adding an item to a Set	61
13.1.2 Removing an item to a Set	61
13.1.3 Union	61
13.1.4 Intersection	61
13.1.5 Set Difference	61
13.1.6 Subset	61
13.2 Operation Analysis	61
13.2.1 TreeSet Vs HashSet Vs Linked Hash Set	61
13.3 Sets and Problem Solving	61
13.3.1 Checking for Uniqueness or Finding Duplicates	61
14 Maps	63
15 Hash Tables	65
15.1 Creating a Hash Function	65
16 Map Reduce	67
16.1 Map	67

V Relationships 69

17 Graphs 71

17.1 Introduction and History	71
17.2 Qualities of a Graph	71
17.2.1 Vertices	71
17.2.2 Edges	71
17.3 Special Graphs and Graph Properties	71
17.3.1 Planar Graphs	71
17.3.2 Bipartite Graphs	72
17.3.3 Directed Acyclic Graphs	72
17.4 Building a Graph	72
17.4.1 Adjacency List	72
17.4.2 Adjacency Matrix	72
17.5 Graph Libraries	72
17.5.1 Java - JUNG	72
17.5.2 Python - networkx	72
17.6 Graphs, Humans, and Networks	72
17.6.1 The Small World	72
17.6.2 Scale Free Graphs	72
17.7 Graphs in Art and Nature - Voronoi Tessellation	72

18 Graph Algorithms 73

18.1 Searching and Traversing	73
18.1.1 Breadth First Search	73
18.1.2 Depth First Search	73
18.2 Shortest Path	73
18.2.1 Dijkstra's Algorithm	73
18.2.2 Bellman-Ford	73
18.3 Topological Sorting	73
18.3.1 Khan's Algorithm	73
18.4 Minimum Spanning Trees	73
18.4.1 Kruskal's Algorithm	73
18.4.2 Prim's Algorithm	73

VI Beyond 75

19 A Nontechnical Introduction to NP-Completeness 77

19.1 The Traveling Salesperson Problem (TSP)	77
19.2 The Longest Path Problem	77
19.3 The Rudrata/Hamiltonian Path Problem	77
19.4 Everything is SAT and 3SAT	77

20 Other Data Structures 79

20.1 Skip Lists	79
---------------------------	----

21 Distributed Hash Tables 81

Part I

Preliminaries

Chapter 1

Introduction

1.1 What is a Data Structures Course

Data Structures is all about defining the different ways we can organize data.

1.2 Why This Book?

This text

1.2.1 Where Does This Book Fit Into a Computer Science Curriculum

Education in Computer Science is based around three core topics: translating the steps of solving a problem into a language a computer can understand, organizing data for solving problems, and techniques that can be used to solve problems. These courses typically covered in a university's introductory course, data structures course, and algorithms course respectively, although different universities decide exactly what content fits in which course. Of course, there is are lot more concepts in computer science, from operating systems and low level programming, to networks and how computers talk to each other. However, all these concepts rely on the knowledge gained in the core courses of programming, data structures, and algorithms.

This textbook is all about Data Structures, the middle section between learning how to program and the more advanced problem solving concepts we learn in Computer Science. Here, we focus on mastering the different ways to organize data, recognize the internal and performative differences between each structure, and learn to recognize the best (if there is one) for a given situation.

1.2.2 What Are My Base Assumptions about the Reader?

This textbook assumes that the student has taken a programming course that has covered the basics. Namely: data types such as ints, doubles, booleans, and strings; if statements, for and while loops; and object orient programming. This book is also suitable for the self taught programmer who has not learned much theoretical programming

1.3 To The Instructor

You'll note that this textbook lacks many of the

Does the lack of varied exercises make cheating on assignments easier as semesters go on? Yes, but that bridge was burned long ago. The cheating student can plagiarize from various websites or anonymously hire another to do their work for them. However, the student who cheats isn't exactly clever and certainly hasn't been exposed to much game theory. They will often cheat from the same source.

1.3.1 How to Use

1.4 To The Student

1.4.1 How to use

Chapter 2

Functions and How They Work

This will be an extremely short chapter, but an important one.

2.1 The Runtime Stack

2.2 Argument vs Parameter

<Programming example?

2.2.1 Does anyone actually care?

I cared enough to look it up. But I had to look it up to double check that I'm correct. In a casual situation or talking with another programmer, I don't actually think anyone would care, but I would take care to get it correct for your assignments and exams, much like you would take care to avoid using "ain't" in a formal essay.

2.3 Passing Arguments

2.3.1 How it Works in Java

2.3.2 How it works in Python

Chapter 3

The Array

3.1 Why Arrays

- because new language: Since this is a data structures course, I assumed that students have had exposure to arrays or array like objects. This chapter goes into a bit of a deeper detail that may have been glossed over and Introduces the topic in the appropriate language if need be.

In other words, I assume you know what an array is , but not necessarily how to use it in Java or Python¹.

- because internal memory lookup
- Because we need to make sure internal knowledge is cohesive (eg arrays of objects are arrays of pointers/references)

3.2 Java and Arrays

The Array is a built in class in Java, but the syntax is a bit unique ²

To create an array in Java we do:

```
Type[] myArray = new Type[sizeOfArray]
```

Here, every item in the array is of whatever **Type** we want, which could be a Class or primitive. Arrays can be whatever integer size we desire, but once set it cannot be changed. This is because to create an array, the computer allocates a contiguous block of memory. If we wanted to resize it, there is no guarantee that this chunk of memory won't have things directly before or after it, preventing us from safely extending its range.

¹Although we use lists in python

²Enough so that I constantly had to look up how to do it my first two years of undergraduate studies, so don't feel too bad if you have to do the same

3.3 Python and Arrays

Python doesn't really do arrays. It instead uses Lists, as we'll see in Chapter 5. `myNotArray3 = []` does not actually make an array like you assume it would in some other language. Instead it makes A list (specifically an arraylist) to contain these items. This works exactly like an array in other languages, but you get some nifty operations that allow us to dynamically resize this array if we need it bigger or smaller.⁴

However, if you really want or need to use an array in python, you can. There are two ways to accomplish this. The first way is the built in `array` package. The python package `numpy` contains

Why would we want to use a

3.3.1 Cool Ways to Build an Array in Python

3.4 How an Array Works

3.4.1 Operations

Retrieving a item stored at an index

Setting the value of an index

3.4.2 Array Internals and the Memory Formula

3.5 Common Array Algorithms

3.5.1 Finding Values in an Array

Finding the Minimum

Finding the Average

3.5.2 Limitations

Most frequent characters Resizing

³On styles: Java convention is to use camel case for variable types (`myVariableName`), while python convention is to use underscores (`my_variable_name`). I will be using the Java style camel-casing for variables throughout the book for consistency and because it is my preference.

⁴We cover the specifics in Chapter ??

Chapter 4

Analyzing Algorithms

4.0.1 Cost

Every function, operation, algorithm, or what have you that a computer performs has a *cost*. In fact, there are always multiples costs; we often just focus on the most important one or two costs. What is most important depends on context.

However, when we measure cost, we need to do abstractly. When we measure the amount of time that an algorithm takes

Time

A time cost is a measure of not just how long it takes a program to finish executing, but also how the length of execution is affected by adding additional item.

Time is almost always *the most important cost*.

Space

Energy

Other costs - Bandwidth

4.1 Big O Notation

- What is big O
- how to read it
- Aside about big omega and theta
- How wrong usage annoys mathematician
- refers to cost in general, but used for time usually
- space complexity
- Common runtimes
- runtimes we'll focus on now

- runtimes we focus on later

4.1.1 Space Complexity

4.2 Examples with Arrays

- Retrieval - refer back to earlier chapter for address lookup
- Replacement
- Linear Search
- Binary Search

4.2.1 Selection Sort

4.2.2 Bubble Sort

4.2.3 Insertion Sort

4.2.4 Other Sorting Algorithms

4.3 The Formal Mathematics of Big O Notation

4.4 Other Notations

4.5 When To Ignore Costs

Part II

Lists

Chapter 5

Array Lists

The first data structure we will be studying is the list. The list is by far the most relatable data structure, as humans deal with lists on a regular basis.

5.1 What is a List?

When you get right down to it, lists are defined by order. We don't have to take advantage of this order, but its there. Populated lists have a first item and they have a last item.

Take a look at this quest below from a hypothetical fantasy game:

Quest: Slay the Dragon of Doom

- Get Sword of Dragonslaying
- Locate the map to Dragon Lair of Doom
- Travel to the Dragon Lair of Doom
- Slay the Dragon of Doom
- Return to the Castle

Here, the order is implied by the contents of the list - you can't beat the dragon without the macguffin and you certainly can't fight it without being able to find it. Generally speaking, going up against a dragon without any preparation is foolhardy in the extreme, but I digress.

Thus, you must get the special sword¹ first, and you must get the map to find the lair before you can physically travel there.

shopping list example

While lists are defined by order, we don't necessarily need to ascribe any meaning to the order. Take a look at the shopping list below:

¹What if its possible to get the map before the sword? We'll see much later this kind of quest and it's requirements are much better handled by a directed acyclic graph in Chapter 17, but this example is fine for teaching lists.

<Shopping List>

While bread is the first item on this list, being the first item in the shopping list in this case has no special meaning. It's not the most important item on the list², nor is it necessarily the item I'm going to pick up first.

Where arrays and lists differ is that lists can grow to an arbitrary size, whereas

- Lists can grow to any arbitrary size, but lists

Data in a list is stored sequentially and array lists do this by creating a an

5.1.1 A note on terminology

An **array list** is a type of list. These are sometimes called dynamic arrays.

As mentioned in Section 3.3, Python doesn't have arrays. If you've been programming in Python, you've been using an array list the entire time you've declared []. They are usually just called lists rather than array lists.

I will be using the Java nomenclature for the majority of the book as this allows me to be clear about the types and implementations of data structures.

5.1.2 Lists in Java

So what does

An Aside about interfaces

Here is the source code

5.2 Generics

5.2.1 What are they?

Before we get to deep into lists, we need to have a discussion about generics. Generics are a way of restricting and specifying what types can go into a collection.

5.2.2 But Why?

5.3 List Operations

5.3.1 Size

We need some easy way of knowing how big our lists are, if for no other reason than to make sure our add and remove methods can figure out their valid indices.

²obviously, that's the cookies

5.3.2 Add

By default, we add items to the end of the list, but we can also add items to any index we want.

When we add an item at some specific index i , the item at i and all indices to the right shift over one. In other words, what was at i is now at $i + 1$, what was at $i + 1$ is at $i + 2$, and so on.

This also an understandable restriction to adding items to a list - we cannot add an item to any index greater than `myList.size() + 1`. Anything greater wouldn't be at the end of the list; it would be beyond it. The same goes for negative indices.

<possible picture showing a legal and illegal add>

5.3.3 Remove

We can remove items from a list much in the same way we can add them. When removing an item at index i , all the indices to the right shift one over to the left.

5.3.4 Get

5.3.5 Set

5.4 ArrayLists

An array list, as you might have guessed, are lists built using *arrays*.³ They work by growing or shrinking the array⁴ automatically as items are added or removed from the list, giving the illusion that the data structure can hold an arbitrary amount of data.

We'll go into the specifics of how this works in Section 5.6.

Java's ArrayLists

Java's `ArrayList`

Python's Lists

Python's lists, such as below:

```
l = [1,2,3] # this is a list, not an array!
```

are actually array lists!

Python uses a different vocabulary for some of the methods we'll be implementing below. For example, take the action of adding an item to a list. Python uses the `append` method to add an item to end of the list and `insert` to put an item into the middle of the list. Java (whose vocabulary we'll be following), uses `add` for both these contexts.

³Shockingly, many of the names we give things at this point actually make sense.

⁴A lie. As you'll see we don't actually change the size of an array; we create a new array of the appropriate size and copy everything over.

5.5 Example Algorithms

```

public static <E> boolean isPermutation(List<E> listA,
    ↪ List<E> listB) {

    if(listA.size() != listB.size()) {
        return false;
    }
    for(int i = 0; i < listA.size() ; i++){
        E item = listA.get(i);
        int countA = 0;
        int countB = 0;

        for (E element : listA) {
            if(item.equals(element)){
                countA++;
            }
        }
        for (E element : listB) {
            if(item.equals(element)){
                countB++;
            }
        }
        if(countA != countB) {
            return false;
        }
    }
    return true;
}

```

5.6 Building an ArrayList

To truly understand how a data structure works we need to implement it ourselves.

5.6.1 Caveats

MyArrayList.java

We will not be implementing the `List` interface. We don't need to implement all the functions to get an understanding of how the fundamentals of an arraylist work. Implementing the list interface would take up a hideous amount of physical space and get in the way of actually understanding the code.

myArrayList.py

For python, this will require some suspension of disbelief, as our array list will require using an array, and as previously discussed, arrays are shirked in favor of arraylists in python. We'll be using a list and pretending it's an array. Silly? Yes. But it will keep our code compact and easier to understand.

5.6.2 Instance Variables

Believe it or not, we only need to keep track of three instance variables to get our arraylist working.

theData We need an array to actually store the items. This is it.

size Size here refers to the total number of items we have stored in the array.

capacity This is the number of items the underlying **array** in our list can hold. It is the maximum size of the list before we have to do something about it. This is not strictly necessary as we can get it by querying **array**'s length. However, making it its own variable will help with the readability.

It is very easy to confuse size and capacity since they both deal with counting how many elements. When I talk about **size**, I am talking about the number of items we have stored in the list we are making. Capacity, on the other hand, depends on the length of the built-in array.

Java

```
public class MyArrayList<E> {

    private E[] theData
    private int size; // how many items are in the list
    private int capacity; // how many items the underlying
        ↪ array can hold
```

First, note the **<E>** after **MyArrayList**. This means that we're saying:

- **MyArrayList** is designed to hold a specific type of object.
- Every **E** we see is a placeholder for some type, which will be that same across the entire lifespan of the object.

Python

```
class MyArrayList(object):
    pass
```

5.6.3 Constructor

We need to set the variables to their initial values upon creating the arraylist. The **size** will be 0, since we won't have any objects stored in it yet. We will set the **initial** capacity to 10. It's a small number and thus won't create much wasted space if we don't fill up **theData**. **theData** will be an empty array of **capacity** length. If **theData** becomes full, we will create a bigger array to hold our items using the **realloc()** method (Section ??)

Java

With our constructor, we have one line of weird black magic in order to create an Array of `E[]`'s.

```
public MyArrayList(){
    size = 0;
    capacity = 10;
    theData = (E[]) new Object[10]; // this generates a
    ↪ warning
}
```

So what's going on with the last line? Typically, when creating an array, we would just say:

```
//doing this in the constructor gives us an error.
TYPE[] myArray = new Type[desired_size];
```

However, Java won't let you create new `E` objects since there's no telling what the constructors will be. This rule extends to arrays of `E`, like so:

```
theData = new E[10];
```

However, when creating a new empty array of objects of any type, we're just making an array of nulls which will eventually be replaced by references to objects. Thus, even though the Java compiler will yell at us about Type safety, we can instead create an array of `Object` and then tell , since all references to any types are the same size.

```
// creating one array of nulls and telling Java
// its another type of array of nulls.
theData = (E[]) new Object[10];
```

Remember how Java and most modern programming languages deal with objects; if you're assigning an object to a variable, like in `Object o = new Object()`, we are storing a reference to that object.

Thus, when we add an item to a list, what really happens is we'll be adding a reference to it - the instructions on how to find it in memory.

Python

Python is fairly straightforward, with the caveat that we are pretending `theData` is an array, and not a list.

```
class MyArrayList(object):
    def __init__(self):
        self.size = 0
        self.capacity = 10
        self.theData = [None]*self.capacity
```

Since built-in lists in Python grow and shrink like we would expect a list to, we initialize `theData` with 10 `None` objects⁵ to mimic the way an array would be initialized.

⁵This is the Python equivalent to the Java `null`.

5.6.4 Size

Now, we will add a size method to our list; fairly straightforward in Java.

```
public int size() {
    return size;
}
```

In Python, we can go ahead and use the built in `__len__` method.

```
def __len__(self):
    return self.size
```

Which can then be invoked with `len(myList)`.

It is worth reinforcing that by keeping track of the size in a variable, retrieving the size of our list is always $O(1)$

5.6.5 Add

In Python, these two add methods are called `append` and `insert` respectively, as python does not support method overloading.

```
public void add(int index, E item) {
    if(index < 0 || index > size) {
        throw new IndexOutOfBoundsException("Index " +
            ↪ index + " out of bounds.");
    }

    if(size == capacity) { // O(n) time...sometimes.
        ↪ Amortized over the cost of adding
        this.reallocate();
    }

    for(int i = size - 1; i >= index; i--) { //If adding to
        ↪ the end... constant
        E temp = theData[i];
        theData[i+1] = temp;
    }

    theData[index] = item;
    size++;
}

def insert(self, index, item):
    self.size += 1
```

Finally, since adding to the end is an extremely common operation, we will overload our `add` method. If our list is provided with only an `item`, as opposed to an `item` and an `index`, we will just add that `item` to the end. Since we already wrote a perfectly good `add` method already that we know works, we'll just have our new method call that one.

```

public boolean add(E item) {
    this.add(size, item); // size is the last valid index
    return true; // What?
}

```

Why are we returning `true` here? The short answer is practice and consistency with future data structures. The long answer is any `Collection` in Java has must have an `add` method and a `List` is type of `Collection`⁶.

`Collection` specifies that `add` must take in an `item` and return a `boolean`. A `true` signals the `add` is successful. A `false` signals that we could not add the `item`. For example, this might happen with a `Set` (Chapter 13)

On the other hand, our Adding at a specific index is unique to lists, and not part of collections, and will always work. Therefore, there's no need to return a `boolean`.

Reallocation

When we need to grow our arraylist

```

private void reallocate(){
    //doubles or 1.5x capacity
    //don't do +1 capacity
    E[] newData = (E[]) new Object[capacity];
    for(int i = 0; i < theData.length; i++) {
        newData[i] = theData[i];
    }

    theData = newData;
    capacity = 2 * capacity;
}

```

We want to double our capacity or at least increase it by 50%, Rather than increasing it by a static number. Consider if we increase the capacity by one each time we reallocated. If we did that, we would have to reallocate every time we added a new item to the list. This would mean that every time we add an item to list, add becomes a linear time - $O(n)$ - operation.

5.6.6 toString and __str__

Now that we supposedly have a method for adding items into the list.

5.6.7 Get and Set

The `get` and `set` methods are fairly straightforward:

get - Given an `index`, retrieve the item stored at that `index`.

⁶Our `MyArrayList` isn't technically a `Collection` since we did not implement the `List` interface, but I digress.

set - Given an **index**, replace the old item stored at that **index** with the provided **item**.

set has one additional quirk, we also want to return the old item we're replacing, just in case the programmer wants to do something with the old item. This would obviate the need for pairing a **get** and **set** call with each other if we want to replace the old item, but do something else with it.

For both **get** and **set**, we want to throw some kind of error if the provided index is out of bounds.

Java

Our **get** is fairly straightforward, but feel free to give more information with the error.

```
public E get(int index) {
    if(index < 0 || index >= size) {
        throw new IndexOutOfBoundsException("Index " +
            ↪ index + " out of bounds.");
    }
    return theData[index];
}
```

The same goes for our **set** method.

```
public E set(int index, E item) {
    if(index < 0 || index >= size) {
        throw new IndexOutOfBoundsException("Index " +
            ↪ index + " out of bounds.");
    }
    E oldItem = theData[index];
    theData[index] = item;

    return oldItem;
}
```

Python

Python supports negative indices. It's up to you whether or not you want to support that, but if you do you need to account for the **capacity** and the **size** of the internal list.

We can take advantage of some of the method calls built into python to make our **myarraylist** support indexing.

5.6.8 Remove

5.7 Analysis

5.8 More Restrictive or Permissive Generics

5.9 A Few More Useful Methods

5.9.1 Constructors

Java's `ArrayList` can optionally take in an integer as an argument. This will start the underlying array's length at that value, rather than the default of 10. This is useful if you know exactly how big your List will be. However, if you aren't removing any items when populating your list, consider using an array instead.

5.9.2 Trim

The `trim()` exists to mitigate the fact that arraylist capacity and size don't often match up.

On the other hand, Python will automatically optimize lists for you CITATION NEEDED.

For other languages, check to see if the equivalent exists

5.9.3 Adding Multiple Items in One Invocation

One common operation is to move or copy all the items from one list to another. In Java, we can use the `addAll()` method, which takes any Java collection as a parameter and all the items in that collection to the object.

```
List<Integer> a = new ArrayList<>();
List<Integer> b = new ArrayList<>();
for(int i = 0; i < 3; i++) {
    a.add(i);
}
for(int i = 3; i < 6; i++) {
    b.add(i);
}
a.addAll(b);
System.out.println(a); // 0 to 5 inclusive
System.out.println(b); // [3, 4, 5]
```

In Python, we can use the `extend()` method on anything that is iterable or use some clever slicing. However, I would always recommend using the method call over the slice, since a method invocation is always more readable.

```
a = [0, 1, 2]
b = [3, 4, 5]
c = a + b # creates a new list, which is not extend
a.extend(b) # adds all of b's items to a
a[len(a):] = b # does the same thing but unreadable.
```

A common beginner mistake in Python is to try to extend a list by calling `append` on the list like so.

```
a = [0, 1, 2]
b = [3, 4, 5]
a.append(b) # a is now [0, 1, 2, [3, 4, 5]]
```

This adds the entire list a single item in the list.

Chapter 6

Linked Lists

Linked lists, also referred to as reference based lists, are the second type of lists typically seen in applications. To be clear a linked list is a list. That means it could be used anywhere an array list can. So Why do we have two objects that are functionally equivalent, two collections that hold things in order, using indexes? The answer is will see, is because each list is good at the thing the other list is less efficient at.

Array based lists use contiguous blocks of memory, allocated all at once and when then capacity of the list is filled up. Utilizing an array makes these types of lists extremely efficient at retrieving an item from a specific index, but adding items anywhere but the end of the list incurs a $O(n)$ runtime.

Linked Lists can do all the things an Array List can, but the underlying structure is completely different. Each item in the list is stored in an Object called a *Node*. Nodes are created as items are added to list, rather than in advance. This means that are not contiguous, but Rather they are scattered throughout the computer's memory. So how in the world do we keep track of where we've stored all these items? The solution resembles the scavenger hunt through the computer's memory. Each node Not only the memory location of the item that is being stored, but the memory location of the next node in the list. An example of this code can be found below¹:

```
// a snippet of the Node Class
// This will live inside the LinkedList class
private static class Node<E> {
    E item;
    Node<E> next;

    public Node(E item) {
        this.item = item;
    }
}
```

¹Why is this class private in Java `private`? An inner class (or private class) is a class that lives within another class. We use this for two reasons: Our nodes only exist to build the linked list, so they don't need to have their own class. The Second reason is What about `static class`? This means that we can create nodes without having to make a Linked List first!

Upon first glance, this code may be very confusing. Each node class contains a reference to a node inside of it. This may give the impression that nodes situated one inside another, like one of those Russian nesting matryoshka dolls. However, keep in mind what the node is actually storing is not other objects, but instead memory locations of where to find them. This means that our linked list is more akin to a scavenger hunt where each objective in the hunt contains the instructions on how to find the next objective.

In other words, the item is the data that is being stored (well actually the memory location, don't forget that), and next refers to the memory location of the next index in the list. Crash course is an excellent video demonstrating this which you can find [here](#):

6.1 Connecting Nodes into a list.

we keep track of only the first and last item in the list, referred to as the head and the tail.

I will be presenting the directions to building a fully functional singly-linked list and doubly-linked list. These directions will differ from the mechanics of how your programming language of choice implements them, but have the same time complexity for their operations. My implementation is constructed with the goal of making the code easy to understand and the decisions that need to be for adding and removing reflect each other. Finally, my code aims to minimize the number of null-pointer exceptions and their ilk a programmer would make.

The full implementations can be found at the end of the Chapter.

6.2 Building a Singly LinkedList

We open up our linked list with a class declaration. If our language uses generics, we specify it there. I'll be choosing not to inherit from the built-in list so we can focus solely on our own code and no external distractions.

In Java, our code begins like this.

```
public class LinkedList<E> { }
```

In Python

```
class LinkedList(object):
    pass
```

6.2.1 The Node

We want the Node class to be a private/internal class, so that the Node we write for a singly linked list and doubly linked list won't get mixed up in our coding environments. This also applies for other data structures that will be using nodes.

```
public class LinkedList<E> {
    private static class Node<E>{
        E item;
```

```

        Node<E> next;

        public Node(E item){
            this.item = item;
        }
    }

class LinkedList(object):
    class Node(object):
        def __init__(self, item) -> None:
            self.item = item
            self.next = None

    pass

```

In the Node private/internal/inner class (and only there), the **this** or **self** refers to the **node** rather than the linked list.

6.2.2 Instance Variables and Constructor

Our linked list `LinkedList` only needs a few Instance variables in order to Function. We need to keep track of the size; Without it we would have no idea what the valid indices are in the list. We need to keep track of the head so we know where to start our scavenger hunt for any particular index or item we're looking for. Finally we'll keep track of the tail. While keeping track of the tail isn't strictly necessary, keeping track of it means that will be able to add an item to the end of the linked list very efficiently ($O(1)$).

The only job of the constructor is to initialize everything to either zero or null.

Finally, it's probably a good idea to go ahead and write getter method for the size of the list.

```

public class LinkedList<E> {
    private Node<E> head;
    private Node<E> tail;
    private int size;

    public int size(){
        return this.size;
    }
}

```

6.2.3 Adding

Our Linked list has two add methods, just like the array list. The first only takes in an item and adds that item to the end of the linked list. It will do this

by calling our second method which takes in an index and an item and inserts that item at that index.²

Let's take a look at our first `add`³ method:

```
public boolean add(E item){
    this.add(this.size, item);
    return true;
}

def add(self, item):
    self.add(self.size, item)
    return True
```

Simple enough! But what about that second `add` method? When we do any kind of operation on a linked list, we need to think about how instance variables in a linked list will be altered. Fortunately, we only have three instance variables: `size`, `head`, and `tail`. When adding to a linked list, the size will always be altered as long as the index is valid. Our list's `head` will only be altered when we add an item to the beginning of the list and our `tail` will only be altered when we add to the end of the list. If the list is empty, then the node for that added item becomes both the head and the tail.

We can simplify our job by breaking the `add` method into five separate cases:

1. The index that we want to add to is out of bounds.
2. We are adding an item to a list that is completely empty. This is going to change the head and tail the list from `null` to something.
3. We are adding an item to index 0, which is going to change the head of the list.
4. We are going to add an item to the end of the list, which means that we are going to change what the tail is.
5. We are adding to some other index in the list, which means that we don't have to bother changing the head or the tail.

Let's start with the first case.

Checking the index is in or out of bounds

Since we passed the check above, we should take a moment before we add an item to address things that need to happen no matter what for Every add condition. Specifically, we need to have a node to hold the item we are adding, and we want to go ahead and increment the size of the list At the end of the method so we don't forget about it.

I will be calling the node that holds the item we are inserting into the list `adding`. As calling it node would be extremely confusing, since we are dealing with so many nodes and other variables like `next` that are also four letters long.

Here's what our changes look like.

²If this sounds familiar, it's because this is precisely what the `add` method in the `arraylist` does. Shocking, right?

³As with the `arraylist`, the `add` method returns a boolean to signify that we were successfully able to add it to the list. This will always be true, but we do this because Java expects this for collections, as explained in `arraylists`

```

public void add(int index, E item) {
    // Scenario 1: index is out of bound
    if(index < 0 || index > size ) { //0(1)
        throw new IndexOutOfBoundsException(index
            ↪ + " is out of bounds");
    }

    Node<E> adding = new Node<E>(item);
    /* the rest of our code*/
    size++;
}

```

Adding to an Empty List

Now let's consider Adding to an empty list. An empty list means the size is 0. If that's the case, we are going to make Adding the new head of the list, As well as the new tail. Just like if you are the only person in line at checkout you are both the first person and the last person in line , this node will also be the first node and the last node in the list , which is why it Will be both the head and tail of the list (at least until we add another item).

```

// Scenario 2: adding to an initially empty list
if(size == 0) {
    head = adding;
    tail = adding;
}

```

Adding an item to the beginning of the list

Adding an item to the beginning of the list means that the node containing it becomes the new head of the list. We do this by attaching Adding to the list, Then informing the list adding is the new head .We do this by setting adding's .next Two point to the current head of the list, then setting The list had to be the node we added.

```

// Scenario 3: adding a new head
else if(index == 0) { //(1)
    adding.next = head;
    head = adding;
}

```

Here, we introduce one of the most important rules we need to follow when working with a linked list : when we are adding an item to the linked list attached the list first , then update the rest of the list to accommodate the new reality.

Failing to do this can have catastrophic results. Consider below Where we set Adding as new head first

```

// Mistakes were made
else if(index == 0) {
    head = adding; // oops
    adding.next = head;
}

```

Note that the number of operations we do here is always the same no matter how big the list is! This means that adding to the head is a constant time operation.

Adding an item to the end of the list

```
// Scenario 4: adding a new tail
else if(index == size ){
    tail.next = adding;
    tail = adding;
}
```

Sidebar: Getting a Node at a Specific Index

```
private Node<E> getNode(int index){ //O(n)
    Node<E> current = head;
    for (int i = 0; i < index; i++) {
        current = current.next;
    }
    return current;
}
```

Inserting an item into a specific index

```
// Scenario 5: everything else
else {
    Node<E> before = getNode(index - 1);
    ↪ //O(n)
    adding.next = before.next;
    before.next = adding;
}
```

The end result

```
public void add(int index, E item) {
    // Scenario 1: index is out of bound
    if(index < 0 || index > size ) { //O(1)
        throw new
            ↪ IndexOutOfBoundsException("Not a valid index :(")
    }

    Node<E> adding = new Node<E>(item);

    // Scenario 2: adding to an initially
    ↪ empty list
    if(size == 0) {
        head = adding;
        tail = adding;
    }
    // Scenario 3: adding a new head
    else if(index == 0) { // O(1)
```

```

        adding.next = head;
        head = adding;
    }
    // Scenario 4: adding a new tail
    else if(index == size ){
        tail.next = adding;
        tail = adding;
    }
    // Scenario 5: everything else
    else {
        Node<E> before = getNode(index
        ↪ -1); //O(n)
        adding.next = before.next;
        before.next = adding;
    }

    size++;
}

```

6.3 Get and Set

Before we got onto our remove method, let's take a look at `get` and `set` very briefly.

6.3.1 Get

Just like with an `ArrayList`, the `get` method returns the item and the specified index. However, since we can't go directly to a specific index like we can with an array or `ArrayList`, we need to iterate thru the `.next` links until we get to the appropriate node. Fortunately, we can just use our `getNode` function that we created when we were writing `add`.

```

public E get(int index) {
    if(index < 0 || index >= size ) {
        throw new
        ↪ IndexOutOfBoundsException(index
        ↪ + " is out of bounds");
    }
    return getNode(index).item;
}

```

6.3.2 Set

`Set` operates very similar to `get`. Remember, `set` also returns the item that is already at the specified index, essentially replacing it.

```

public E set(int index, E item) {
    if(index < 0 || index >= size ) { //O(1)

```

```

        throw new
        ↪ IndexOutOfBoundsException(index
        ↪ + " is out of bounds");
    }
    Node<E> node = getNode(index);
    E toReturn = node.item;
    node.item = item;

    return toReturn;
}

```

6.4 Remove

6.5 Analysis

Array lists and linked lists are both extremely powerful objects that fulfill the same purpose, but in radically different ways.

6.5.1 Some Algorithms Play Better

6.6 Potential Project/Practice/Labs

6.7 Source Code

```

from typing import Generic, TypeVar

E = TypeVar('E')

class LinkedList(Generic[E]):

    class Node(Generic[E]):
        def __init__(self, item: E) -> None:
            self.item = item
            self.next = None

    def __init__(self) -> None:
        self.head = None
        self.tail = None
        self.size = 0

    def __len__(self) -> int:
        return self.size

    def getNode(self, index: int) -> Node:
        current = self.head
        for i in range(index):
            current = current.next

```



```

        return current

def add(self, item: E) -> bool:
    self.add(index, index, item)
    return True

def add(self, index: int, item: E) -> None:
    if(index < 0 or index > self.size):
        raise Exception("Invalid add at index " + str(index)
            ↪ " with item" + str(item) + ".")

    adding = self.Node(item)
    if(self.size == 0):
        self.head = adding
        self.tail = adding
    elif(index == 0):
        adding.next = self.head
        self.head = adding
    elif(index == self.size):
        self.tail.next = adding
        self.tail = adding
    else:
        before = self.getNode(index - 1)
        adding.next = before.next
        before.next = adding

    self.size += 1

def remove(self, index: int) -> E:
    if(index < 0 or index >= self.size):
        raise Exception("Invalid remove at index " +
            ↪ str(index) + ".")

    toReturn = None
    if self.size == 1:
        self.head = None
        self.tail = None
    elif index == 0:
        toReturn = self.head.item
        self.head = self.head.next

    self.size -= 1
    return toReturn

l = LinkedList()
print(len(l))

```


Chapter 7

Stacks

Our next data structure is the Stack. The stack may seem unnecessary as a data structure after we introduce its features. After all, can't a list do all the things that a stack can do and more?

Working with the limited operations of a allows us to approach problems with a different mindset.

7.1 Stack Operations

7.2 Building a Stack

7.3 Built-in Stacks

7.4 Solving Problems with A Stack

7.5 Mazes - Stacks and Backtracking

7.6 Discrete Finite Automata

Chapter 8

Queues

A Queue (pronounced by saying the first letter and ignoring all the others) is a data structure which emulates the real world functionality of standing in a line (or queue, for those from Commonwealth nations). In a Queue, items are processed in the order they are inserted into the Queue. So if Alice enters the Queue, followed by Bob, followed by Carla, Alice would be the first to leave the Queue, then Bob, and then Carla.

The use cases for Queues are fairly obvious

8.1 Linked Based Implementation

8.2 Array Based Implementation

We could use

Part III

Recursion

Chapter 9

Recursion

9.1 Introduction

9.2 Recursive Mathematics

9.2.1 Fibonacci

As it turns out, while this technically works...it's pretty terrible. In short, using recursion, I managed to accidentally¹ write an $O(2^n)$, or exponential time, algorithm. This is very bad. This means increasing n by one *doubles* the runtime of our algorithm!

9.3 Redoing Things With Recursion

Many of the things we are about to see should not be actually used and serve only as examples, like our `printThis` function

9.3.1 Printing Recursively

9.3.2 Recursive Linear Search

9.4 Binary Search

9.4.1 Runtime Analysis

How to not be scared of logarithms

You may have learned that logarithms are the inverse operation to exponentiation.

This is an utterly useless definition when programming.

A more way of thinking about logarithms is "how many times can I recursively split something?" For example, $\log_b x$ asks "how many times can I recursively split my x items into b separate piles?"

¹All right, I did this totally on purpose.

A more concrete example: $\log_2 16 = 4$, not because $2^4 = 16$, but because a pile of 16 items can be split in half into two piles of 8, each pile of 8 can be split in half into two piles of 4, the 4's can be split into 2's, the 2's into 1's — four splits total:

<picture>

Back to it.

9.5 Recursive Backtracking

Recursion really comes in handy when we are trying to solve complex puzzles. One of the most famous examples of

The Recursive Backtracking Algorithm

9.5.1 Mazes Again

9.5.2 The Eight Queens Puzzle

Brute Force Solution

Recursive Solution Outline

A Place Holder For Validity

Performing the Recursion

Checking just One condition

Checking all the Conditions

9.5.3 Additional Problems left to the Reader

Knight's Tour

Sudoku

9.6 Recursive Combinations

9.7 Recursion and Puzzles

9.8 Recursion and Art

9.9 Recursion and Nature

Chapter 10

Trees

Our next major data structure is trees. Specifically, we will be looking at binary search trees.

Trees are an excellent data structure for storing things since they implement all the operations we care about for collections in logarithmic time¹

However, trees are not without limitations. Trees will only work with data that can be stored hierarchically or in an order.

10.1 The Parts of a Tree

The first thing we need to do when introducing trees is define a vocabulary.

Much like the linked list, a tree is made of nodes. However, unlike a linked list, nodes in a tree are not arranged in a line. Instead, they are arranged in a hierarchy.

Each node sits above multiple other nodes, with the nodes below it being referred to as their children or child nodes. The node connecting all these children is called the parent.

<A picture of one node, Represented by a circle with four arrows coming out below it. Each arrow points to yet another node. The Node with the arrows coming out of it is the parent, and the nodes below it are the children >

This relationship can be extended Ad infinitum as we can see with the picture below

<Picture with nodes labeled>

However anything above grandchild and grandparent just becomes tedious, so we tend to Generalize this relationship to ancestors and descendants. A key point here is to remember that while we are borrowing terms from the family tree, nodes will only have one parent. Each node can have multiple children, however.

We refer to the links connect each of the nodes as branches or links or edges. This tends to be a matter of personal preference.

¹Specifically, Trees implement everything in average case logarithmic time and worst case linear time, but if we do a bit of extra work and make it a self balancing binary tree (which will seem much later in this chapter) we can make this tree worst case logarithmic for all operations

Finally, we have one special node that sits above all the other nodes. This node is the root and it is analogous to the head of a linked list. All of our operations will start at the root of the tree².

Remember, programmers are stereotypically outdoors of averse, so they may have forgotten what a real tree looks like. Thus, we'll see that the root of the tree is at the top of the tree and our leaves are at the bottom⁴.

10.1.1 Where the Recursion comes in

There is a reason we learned recursion before we introduce trees. Trees are the exemplar recursive data structure.

Each tree has a root and that root has children. If we view each of those children as the root of their own subtree, this can make our algorithms for adding, removing, and searching extremely easy to write.

<picture Of tree, the recursive subtrees are dash circled.>

<Picture of the left subtree, with its trees circled>

10.2 Binary Search Trees

A diagram of a binary search tree. It is made up of nodes, represented by circles, and edges (also called links or branches), represented by arrows.

5

10.3 Building a Binary Search Tree

10.3.1 The Code Outline

We use the `Comparable` class in Java to require that all objects stored in the tree have a **total ordering**⁶. In practice, this means that anything `Comparable` can be sorted.

Python, of course, doesn't need these restrictions.

```
public class BinaryTree<E> extends Comparable<E> {
}
```

Much like our Linked List, we don't need much in the way of instance variables. We'll create a `root` to keep track of the starting place for our tree and size to keep track of how many items we have stored.

Finally, we will also create our inner `Node` class for the Tree. It needs to hold the item and the locations of the left and right children. We'll also go ahead and add a constructor and a method for printing out the item in the node (`toString` in Java and `__str__` in Python).

²Remember, programmers are stereotypically outdoors of averse, so they may have forgotten what a real tree looks like. Thus, we'll see that the root of the tree is at the top of the tree and our leaves are at the bottom³.

⁴Or maybe it's some weird hydroponic zero-G kind of thing.

⁵An aside about array based implementations.

⁶The formal definition is as follows

```
public class BinaryTree<E extends Comparable<E>> {
    private Node<E> root;
    private int size;

    public BinaryTree() {
        this.root = null;
    }

    /* Other code will go here.*/

    private static class Node<E extends
↪ Comparable<E>> {
        private E item;
        private Node<E> left; // left child
        private Node<E> right; // right child
        public Node(E item) {
            this.item = item;
        }
        public String toString() {
            return item.toString();
        }
    }
}
```

10.3.2 Add

10.3.3 Contains

10.3.4 Delete

Chapter 11

Heaps

11.1 Priority Queues

11.2 Removing From other locations

Chapter 12

Sorting

12.1 Quadratic-Time Algorithms

12.1.1 Bubble Sort

12.1.2 Selection Sort

12.1.3 Insertion Sort

12.2 Log-Linear Sorting Algorithms

The most commonly used sorting algorithms take $O(n \lg(n))$ time. This is the hard limit on runtime

12.2.1 Tree Sort

The tree sort is the simplest algorithm to we will cover. Performing Tree sort is a matter of three simple steps

1. Create a tree.
2. Load the items you want to sort into the tree.
3. Perform an inorder traversal of the tree.

The performance of this algorithm depends completely on the type of tree we create for this algorithm. Using a self-balancing binary search tree, adding n items to the tree takes $O(n \lg(n))$ and an in order traversal takes $O(n)$ steps, for a grand total of $O(n)$ runtime. Using a binary search tree that does not self balance means that there is a worst case scenario of $O(n^2)$ for adding all the n items.

Using a tree also means we use extra space since all the data has to be moved into a tree, using $O(n)$ space.

12.2.2 Heap Sort

You might expect that heapsort deserves the same treatment as treesort. After all, a heap has the same structure as a tree and both are constructed to perform operations in $\log n$ time.

12.2.3 Heapify**12.2.4 Quick Sort****12.2.5 Merge Sort****12.3 Unique Sorting Algorithms****12.3.1 Shell Sort****12.3.2 Radix Sort****12.4 State of the Art Sorting Algorithms****12.4.1 Tim Sort****12.4.2 Quick Sort****12.5 But What if We Add More Computers: Parallelization and Distributed Algorithms**

Parallel VS Distributed

12.6 Further Reading**12.6.1 Pedagogical Sorting Algorithms**

Sleep Sort

Part IV

Hashing

Chapter 13

Sets

Sets programmed implementations of mathematical sets

13.1 Operations

We will use Venn diagrams to graphically demonstrate operates with two sets

13.1.1 Adding an item to a Set

13.1.2 Removing an item to a Set

13.1.3 Union

13.1.4 Intersection

13.1.5 Set Difference

13.1.6 Subset

13.2 Operation Analysis

Most sets are implemented using a Hash Table.

13.2.1 TreeSet Vs HashSet Vs Linked Hash Set

13.3 Sets and Problem Solving

Sets are super efficient checklists.

13.3.1 Checking for Uniqueness or Finding Duplicates

Chapter 14

Maps

Chapter 15

Hash Tables

15.1 Creating a Hash Function

Chapter 16

Map Reduce

16.1 Map

The `map()` operation¹ is a powerful function that may require us to think differently about the way we have approached programming so far.

The map operation takes in 2 arguments, a collection and a function to apply to every item in the collection

When we are writing functions , we are creating new verbs for our programming language to use . These verbs take in arguments, nouns that we may have declared or defined ourselves. But one thing that we May not have done yet is passing a function as an argument to another function.

This is not an uncommon operation in mathematics Example listed below

The semantics of this in every programming language is different , but the concept is the same

Why introduces here? Because a lot of common operations that can be done with map reduce involve using hash tables

¹It is mildly confusing that there is a `map` data structure and a `map()` operation, so I will be marking the `map()` operation with a function invocation.

Part V

Relationships

Chapter 17

Graphs

In some ways, Graphs are the most important data structure. Graphs represent and model relationships, and humans are defined by relationships. The archetypical examples of graphs used to be maps and the distances between landmarks or looking for the shortest path.

With the advent of social media, we can talk about graphs with a few examples that might be easier to intuit.

17.1 Introduction and History

17.2 Qualities of a Graph

The physical layout of a graph doesn't actually matter¹

17.2.1 Vertices

- Vertices must be unique.

17.2.2 Edges

Undirected Edges

Directed Edges

Weighted Edges

17.3 Special Graphs and Graph Properties

17.3.1 Planar Graphs

Graphs that are planar can have their vertices and edges laid out in such a way that no two edges will cross.

¹Some properties, such as whether a graph is *planar* or *bipartite* effectively care if a graph can be physically laid out in a certain way.



Figure 17.1: The wings of a dragonfly. Credit: Joi Ito (CC BY 2.0)

17.3.2 Bipartite Graphs

17.3.3 Directed Acyclic Graphs

17.4 Building a Graph

17.4.1 Adjacency List

17.4.2 Adjacency Matrix

Matrix multiplication and GPU Abuse

17.5 Graph Libraries

17.5.1 Java - JUNG

17.5.2 Python - networkx

There is only one realistic choice for using graphs in Python. The package `networkx` is extremely powerful, extremely versatile, and actively maintained.

17.6 Graphs, Humans, and Networks

17.6.1 The Small World

The Milgram Experiment

The Less-Known Milgram Experiment

17.6.2 Scale Free Graphs

17.7 Graphs in Art and Nature - Voronoi Tessellation

Chapter 18

Graph Algorithms

18.1 Searching and Traversing

18.1.1 Breadth First Search

18.1.2 Depth First Search

18.2 Shortest Path

18.2.1 Dijkstra's Algorithm

Improving The Algorithm

Failure Cases

18.2.2 Bellman-Ford

18.3 Topological Sorting

18.3.1 Khan's Algorithm

18.4 Minimum Spanning Trees

18.4.1 Kruskal's Algorithm

18.4.2 Prim's Algorithm

Part VI

Beyond

Chapter 19

A Nontechnical Introduction to NP-Completeness

19.1 The Traveling Salesperson Problem (TSP)

19.2 The Longest Path Problem

19.3 The Rudrata/Hamiltonian Path Problem

19.4 Everything is SAT and 3SAT

Chapter 20

Other Data Structures

20.1 Skip Lists

Chapter 21

Distributed Hash Tables